

NAME: YASH AWARE

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter)
# will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# importing libraries and magic functions

import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score

from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
%config InlineBackend.figure_format ='retina'
%matplotlib inline

os.chdir("C:\\\\Users\\\\PC-6\\\\Desktop")
# read data
df = pd.read_csv('framingham.csv')

# first glimpse at data
df.head(20)

# data shape
df.shape
```

## # data types

	male	age	education	currentSmoker	cigsPerDay	BPMeds	
0	1	39	4.0	0	0.0	0.0	
1	0	46	2.0	0	0.0	0.0	
2	1	48	1.0	1	20.0	0.0	
3	0	61	3.0	1	30.0	0.0	
4	0	46	3.0	1	23.0	0.0	
5	0	43	2.0	0	0.0	0.0	
6	0	63	1.0	0	0.0	0.0	
7	0	45	2.0	1	20.0	0.0	
8	1	52	1.0	0	0.0	0.0	
9	1	43	1.0	1	30.0	0.0	
10	0	50	1.0	0	0.0	0.0	
11	0	43	2.0	0	0.0	0.0	
12	1	46	1.0	1	15.0	0.0	
13	0	41	3.0	0	0.0	1.0	
14	0	39	2.0	1	9.0	0.0	
15	0	38	2.0	1	20.0	0.0	
16	1	48	3.0	1	10.0	0.0	
17	0	46	2.0	1	20.0	0.0	
18	0	38	2.0	1	5.0	0.0	
19	1	41	2.0	0	0.0	0.0	
0							
glucose	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate
0	0	0	195.0	106.0	70.0	26.97	80.0
77.0							

1	0	0	250.0	121.0	81.0	28.73	95.0	
76.0								
2	0	0	245.0	127.5	80.0	25.34	75.0	
70.0								
3	1	0	225.0	150.0	95.0	28.58	65.0	
103.0								
4	0	0	285.0	130.0	84.0	23.10	85.0	
85.0								
5	1	0	228.0	180.0	110.0	30.30	77.0	
99.0								
6	0	0	205.0	138.0	71.0	33.11	60.0	
85.0								
7	0	0	313.0	100.0	71.0	21.68	79.0	
78.0								
8	1	0	260.0	141.5	89.0	26.36	76.0	
79.0								
9	1	0	225.0	162.0	107.0	23.61	93.0	
88.0								
10	0	0	254.0	133.0	76.0	22.91	75.0	
76.0								
11	0	0	247.0	131.0	88.0	27.64	72.0	
61.0								
12	1	0	294.0	142.0	94.0	26.31	98.0	
64.0								
13	1	0	332.0	124.0	88.0	31.31	65.0	
84.0								
14	0	0	226.0	114.0	64.0	22.35	85.0	
NaN								
15	1	0	221.0	140.0	90.0	21.35	95.0	
70.0								
16	1	0	232.0	138.0	90.0	22.37	64.0	
72.0								
17	0	0	291.0	112.0	78.0	23.38	80.0	
89.0								
18	0	0	195.0	122.0	84.5	23.24	75.0	
78.0								
19	0	0	195.0	139.0	88.0	26.88	85.0	
65.0								
TenYearCHD								
0	0							
1	0							
2	0							
3	1							
4	0							
5	0							
6	1							
7	0							
8	0							

```

9          0
10         0
11         0
12         0
13         0
14         0
15         1
16         0
17         1
18         0
19         0

(4240, 16)

# check for duplicates
duplicate_df = df[df.duplicated()]
duplicate_df

Empty DataFrame
Columns: [male, age, education, currentSmoker, cigsPerDay, BPMeds,
prevalentStroke, prevalentHyp, diabetes, totChol, sysBP, diaBP, BMI,
heartRate, glucose, TenYearCHD]
Index: []

# checking for missing values
df.isna().sum()
null = df[df.isna().any(axis=1)]
null

male          0
age           0
education     105
currentSmoker 0
cigsPerDay    29
BPMeds        53
prevalentStroke 0
prevalentHyp   0
diabetes       0
totChol        50
sysBP          0
diaBP          0
BMI            19
heartRate      1
glucose        388
TenYearCHD     0
dtype: int64

      male  age  education  currentSmoker  cigsPerDay  BPMeds \
14      0   39       2.0           1       9.0     0.0
21      0   43       1.0           0       0.0     0.0
26      0   60       1.0           0       0.0     0.0

```


33	1	61	NaN	1	5.0	0.0	
36	1	56	NaN	0	0.0	0.0	
...	...	...	...	...	...	...	
4208	0	51	1.0	1	9.0	0.0	
4229	0	51	3.0	1	20.0	0.0	
4230	0	56	1.0	1	3.0	0.0	
4235	0	48	2.0	1	20.0	NaN	
4236	0	44	1.0	1	15.0	0.0	
BMI		prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP
14		0	0	0	226.0	114.0	64.0
22.35		0	0	0	185.0	123.5	77.5
21		0	0	0	260.0	110.0	72.5
29.89		0	0	0	175.0	134.0	82.5
26		0	0	0	257.0	153.5	102.0
26.59		0	0	0	340.0	152.0	76.0
33		0	0	0	251.0	140.0	80.0
18.59		0	0	0	268.0	170.0	102.0
36		0	0	0	248.0	131.0	72.0
28.09		0	0	0	210.0	126.5	87.0
...		...	...	...	...	...	...
4208		0	0	0	340.0	152.0	76.0
25.74		0	1	0	251.0	140.0	80.0
4229		0	1	0	268.0	170.0	102.0
25.60		0	0	0	248.0	131.0	72.0
4230		0	0	0	210.0	126.5	87.0
22.89		0	0	0	340.0	152.0	76.0
4235		0	0	0	251.0	140.0	80.0
22.00		0	0	0	268.0	170.0	102.0
4236		0	0	0	248.0	131.0	72.0
19.16		0	0	0	210.0	126.5	87.0
	heartRate	glucose	TenYearCHD				
14	85.0	NaN	0				
21	70.0	NaN	0				
26	65.0	NaN	0				
33	72.0	75.0	1				
36	72.0	75.0	0				
...	...	...	...				
4208	70.0	NaN	0				
4229	75.0	NaN	0				
4230	57.0	NaN	0				
4235	84.0	86.0	0				
4236	86.0	NaN	0				

[582 rows x 16 columns]


```
import matplotlib.pyplot as plt

# Assuming df is your dataframe
fig = plt.figure(figsize=(15, 20))
df.hist() # Automatically handles subplot layout
plt.tight_layout()
plt.show()

array([[<Axes: title={'center': 'male'}>,
       <Axes: title={'center': 'age'}>,
       <Axes: title={'center': 'education'}>,
       <Axes: title={'center': 'currentSmoker'}>],
      [<Axes: title={'center': 'cigsPerDay'}>,
       <Axes: title={'center': 'BPMeds'}>,
       <Axes: title={'center': 'prevalentStroke'}>,
       <Axes: title={'center': 'prevalentHyp'}>],
      [<Axes: title={'center': 'diabetes'}>,
       <Axes: title={'center': 'totChol'}>,
       <Axes: title={'center': 'sysBP'}>,
       <Axes: title={'center': 'diaBP'}>],
      [<Axes: title={'center': 'BMI'}>,
       <Axes: title={'center': 'heartRate'}>,
       <Axes: title={'center': 'glucose'}>,
       <Axes: title={'center': 'TenYearCHD'}>]], dtype=object)
```



<Figure size 1500x2000 with 0 Axes>



```


import seaborn as sns
import matplotlib.pyplot as plt

# Assuming df is your dataframe
df_corr = df.corr()

# Plot heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(df_corr, annot=True, cmap='coolwarm', fmt='.2f',
cbar=True)
plt.title('Correlation Matrix')
plt.show()

<Figure size 1200x800 with 0 Axes>
<Axes: >
Text(0.5, 1.0, 'Correlation Matrix')

```



```
# Dropping columns education and glucose
df = df.drop(['education'], axis=1)

# Checking for more missing data
df.isna().sum()
```

male	0
age	0
currentSmoker	0
cigsPerDay	29
BPMed	53
prevailStroke	0
prevailHyp	0
diabetes	0
totChol	50
sysBP	0
diaBP	0
BMI	19
heartRate	1
glucose	388

```

TenYearCHD          0
dtype: int64

# Dropping all rows with missing data
df = df.dropna()
df.isna().sum()
df.columns

male              0
age               0
currentSmoker    0
cigsPerDay       0
BPMeds            0
prevalentStroke  0
prevalentHyp     0
diabetes          0
totChol           0
sysBP             0
diaBP             0
BMI               0
heartRate         0
glucose           0
TenYearCHD        0
dtype: int64

Index(['male', 'age', 'currentSmoker', 'cigsPerDay', 'BPMeds',
       'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol',
       'sysBP',
       'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD'],
      dtype='object')

# Identify the features with the most importance for the outcome variable Heart Disease

from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2

# separate independent & dependent variables
X = df.iloc[:,0:14] #independent columns
y = df.iloc[:, -1]   #target column i.e price range

# apply SelectKBest class to extract top 10 best features
bestfeatures = SelectKBest(score_func=chi2, k=10)
fit = bestfeatures.fit(X,y)
dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(X.columns)

#concat two dataframes for better visualization
featureScores = pd.concat([dfcolumns,dfscores],axis=1)
featureScores.columns = ['Specs','Score'] #naming the dataframe

```

```

columns
print(featureScores.nlargest(11, 'Score')) #print 10 best features

          Specs      Score
9        sysBP  667.109932
13       glucose  402.409837
1         age   297.974980
8        totChol 252.958627
3       cigsPerDay 185.115196
10       diaBP  142.920436
6    prevalentHyp  82.342164
7       diabetes  31.711253
4        BPMeds  26.116583
0         male   19.178560
11        BMI   17.108210

featureScores = featureScores.sort_values(by='Score', ascending=False)
featureScores

          Specs      Score
9        sysBP  667.109932
13       glucose  402.409837
1         age   297.974980
8        totChol 252.958627
3       cigsPerDay 185.115196
10       diaBP  142.920436
6    prevalentHyp  82.342164
7       diabetes  31.711253
4        BPMeds  26.116583
0         male   19.178560
11        BMI   17.108210
5  prevalentStroke  8.480982
12     heartRate  3.635480
2  currentSmoker  0.904429

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(20,5))

# Assuming featureScores is a DataFrame
sns.barplot(x='Specs', y='Score', data=featureScores)

# Rotate the x-axis labels to prevent overlap
plt.xticks(rotation=45, ha='right', fontsize=12)

# Remove the warning by removing the palette argument
plt.box(False)
plt.title('Feature importance', fontsize=16)
plt.xlabel('\n Features', fontsize=14)
plt.ylabel('Importance \n', fontsize=14)

```

```
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)

plt.show()

<Figure size 2000x500 with 0 Axes>

<Axes: xlabel='Specs', ylabel='Score'>

([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13],
 [Text(0, 0, 'sysBP'),
  Text(1, 0, 'glucose'),
  Text(2, 0, 'age'),
  Text(3, 0, 'totChol'),
  Text(4, 0, 'cigsPerDay'),
  Text(5, 0, 'diaBP'),
  Text(6, 0, 'prevalentHyp'),
  Text(7, 0, 'diabetes'),
  Text(8, 0, 'BPMeds'),
  Text(9, 0, 'male'),
  Text(10, 0, 'BMI'),
  Text(11, 0, 'prevalentStroke'),
  Text(12, 0, 'heartRate'),
  Text(13, 0, 'currentSmoker')])

Text(0.5, 1.0, 'Feature importance')

Text(0.5, 0, '\n Features')


Text(0, 0.5, 'Importance \n')

([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13],
 [Text(0, 0, 'sysBP'),
  Text(1, 0, 'glucose'),
  Text(2, 0, 'age'),
  Text(3, 0, 'totChol'),
  Text(4, 0, 'cigsPerDay'),
  Text(5, 0, 'diaBP'),
  Text(6, 0, 'prevalentHyp'),
  Text(7, 0, 'diabetes'),
  Text(8, 0, 'BPMeds'),
  Text(9, 0, 'male'),
  Text(10, 0, 'BMI'),
  Text(11, 0, 'prevalentStroke'),
  Text(12, 0, 'heartRate'),
  Text(13, 0, 'currentSmoker')])

(array([ 0., 100., 200., 300., 400., 500., 600., 700., 800.]),
 [Text(0, 0.0, '0'),
  Text(0, 100.0, '100'),
  Text(0, 200.0, '200'),
```

```

Text(0, 300.0, '300'),
Text(0, 400.0, '400'),
Text(0, 500.0, '500'),
Text(0, 600.0, '600'),
Text(0, 700.0, '700'),
Text(0, 800.0, '800')))
```



```

# selecting the 10 most impactful features for the target variable
features_list = featureScores["Specs"].tolist()[:10]
features_list

['sysBP',
 'glucose',
 'age',
 'totChol',
 'cigsPerDay',
 'diaBP',
 'prevalentHyp',
 'diabetes',
 'BPMeds',
 'male']

# Create new dataframe with selected features

df = df[['sysBP',
 'glucose', 'age', 'totChol', 'cigsPerDay', 'diaBP', 'prevalentHyp', 'diabetes',
 'BPMeds', 'male', 'TenYearCHD']]
df.head()

   sysBP  glucose    age  totChol  cigsPerDay  diaBP  prevalentHyp
diabetes \
0  106.0      77.0    39    195.0        0.0    70.0                 0
0
1  121.0      76.0    46    250.0        0.0    81.0                 0
0
2  127.5      70.0    48    245.0       20.0    80.0                 0
```

```

0
3 150.0      103.0    61      225.0          30.0    95.0          1
0
4 130.0      85.0     46      285.0          23.0    84.0          0
0

    BPMeds   male  TenYearCHD
0      0.0      1          0
1      0.0      0          0
2      0.0      1          0
3      0.0      0          1
4      0.0      0          0

import matplotlib.pyplot as plt
import seaborn as sns

# Calculate correlation matrix
df_corr = df.corr()

# Plot heatmap
plt.figure(figsize=(12, 8)) # Adjust figure size for better readability
sns.heatmap(df_corr, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=0.5, cbar_kws={'shrink': 0.8})

# Optional: Title and labels
plt.title('Correlation Heatmap', fontsize=16)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)

plt.show()


<Figure size 1200x800 with 0 Axes>

<Axes: >

Text(0.5, 1.0, 'Correlation Heatmap')

(array([ 0.5,  1.5,  2.5,  3.5,  4.5,  5.5,  6.5,  7.5,  8.5,  9.5,
10.5]),
 [Text(0.5, 0, 'sysBP'),
  Text(1.5, 0, 'glucose'),
  Text(2.5, 0, 'age'),
  Text(3.5, 0, 'totChol'),
  Text(4.5, 0, 'cigsPerDay'),
  Text(5.5, 0, 'diaBP'),
  Text(6.5, 0, 'prevalentHyp'),
  Text(7.5, 0, 'diabetes'),
  Text(8.5, 0, 'BPMeds'),
  Text(9.5, 0, 'male'),
  Text(10.5, 0, 'TenYearCHD')])
```

```
(array([ 0.5,  1.5,  2.5,  3.5,  4.5,  5.5,  6.5,  7.5,  8.5,  9.5,
10.5]),
[Text(0, 0.5, 'sysBP'),
Text(0, 1.5, 'glucose'),
Text(0, 2.5, 'age'),
Text(0, 3.5, 'totChol'),
Text(0, 4.5, 'cigsPerDay'),
Text(0, 5.5, 'diaBP'),
Text(0, 6.5, 'prevalentHyp'),
Text(0, 7.5, 'diabetes'),
Text(0, 8.5, 'BPMeds'),
Text(0, 9.5, 'male'),
Text(0, 10.5, 'TenYearCHD')])
```



```
# Checking for outliers
df.describe()
sns.pairplot(df)
```

	sysBP	glucose	age	totChol	cigsPerDay
\count	3751.000000	3751.000000	3751.000000	3751.000000	3751.000000
mean	132.368435	81.880032	49.573447	236.928019	9.008531
std	22.046522	23.882233	8.570204	44.611594	11.925097
min	83.500000	40.000000	32.000000	113.000000	0.000000
25%	117.000000	71.000000	42.000000	206.000000	0.000000
50%	128.000000	78.000000	49.000000	234.000000	0.000000
75%	144.000000	87.000000	56.000000	264.000000	20.000000
max	295.000000	394.000000	70.000000	696.000000	70.000000
	diaBP	prevalentHyp	diabetes	BPMeds	
male \count	3751.000000	3751.000000	3751.000000	3751.000000	3751.000000
mean	82.938550	0.311917	0.027193	0.030392	0.445215
std	11.932779	0.463338	0.162666	0.171686	0.497056
min	48.000000	0.000000	0.000000	0.000000	0.000000
25%	75.000000	0.000000	0.000000	0.000000	0.000000
50%	82.000000	0.000000	0.000000	0.000000	0.000000
75%	90.000000	1.000000	0.000000	0.000000	1.000000
max	142.500000	1.000000	1.000000	1.000000	1.000000
	TenYearCHD				
count	3751.000000				
mean	0.152493				
std	0.359546				
min	0.000000				
25%	0.000000				
50%	0.000000				
75%	0.000000				
max	1.000000				
<seaborn.axisgrid.PairGrid at 0x239006469f0>					

```

import seaborn as sns
import matplotlib.pyplot as plt

# Describe the dataset for summary statistics
print(df.describe())

# Pairplot to visualize pairwise relationships
sns.pairplot(df)

# Optional: Show the plot
plt.show()
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1

# Detecting outliers
outliers = ((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).sum()
print("Outliers detected:\n", outliers)

```

	sysBP	glucose	age	totChol	cigsPerDay
count	3751.000000	3751.000000	3751.000000	3751.000000	3751.000000
mean	132.368435	81.880032	49.573447	236.928019	9.008531
std	22.046522	23.882233	8.570204	44.611594	11.925097
min	83.500000	40.000000	32.000000	113.000000	0.000000
25%	117.000000	71.000000	42.000000	206.000000	0.000000
50%	128.000000	78.000000	49.000000	234.000000	0.000000
75%	144.000000	87.000000	56.000000	264.000000	20.000000
max	295.000000	394.000000	70.000000	696.000000	70.000000


  

	diaBP	prevalentHyp	diabetes	BPMeds
male \				
count	3751.000000	3751.000000	3751.000000	3751.000000
3751.000000				
mean	82.938550	0.311917	0.027193	0.030392
0.445215				
std	11.932779	0.463338	0.162666	0.171686
0.497056				
min	48.000000	0.000000	0.000000	0.000000
0.000000				
25%	75.000000	0.000000	0.000000	0.000000
0.000000				
50%	82.000000	0.000000	0.000000	0.000000

```
0.000000
75%      90.000000      1.000000      0.000000      0.000000
1.000000
max      142.500000      1.000000      1.000000      1.000000
1.000000

      TenYearCHD
count  3751.000000
mean    0.152493
std     0.359546
min    0.000000
25%    0.000000
50%    0.000000
75%    0.000000
max    1.000000
```

```
<seaborn.axisgrid.PairGrid at 0x23918af2240>
```



### Outliers detected:

sysBP	110
glucose	181
age	0
totChol	42
cigsPerDay	11
diaBP	69
prevailentHyp	0
diabetes	102
BPMed	114
male	0

```

TenYearCHD      572
dtype: int64

import seaborn as sns
import matplotlib.pyplot as plt


# Create a boxplot to visualize the distribution of 'totChol'
sns.boxplot(df['totChol'])
plt.title('Boxplot of Total Cholesterol')
plt.show()

# Identify outliers in the 'totChol' column (values greater than 500
# in this case)
outliers = df[df['totChol'] > 500]

# Display the outliers
print(outliers)

<Axes: ylabel='totChol'>
Text(0.5, 1.0, 'Boxplot of Total Cholesterol')

```



	sysBP	glucose	age	totChol	cigsPerDay	diaBP	prevalentHyp
diabetes \	1111	159.5	140.0	52	600.0	0.0	94.0
							1

```

1
3160 157.0      84.0   51     696.0          9.0    87.0          1
0

    BPMeds  male  TenYearCHD
1111      0.0     0         1
3160      0.0     1         0


import seaborn as sns
import matplotlib.pyplot as plt

# Drop outliers where 'totChol' is greater than 599
df = df.drop(df[df['totChol'] > 599].index)

# Create a boxplot to visualize the distribution of 'totChol' after
# dropping outliers
sns.boxplot(df['totChol'])
plt.title('Boxplot of Total Cholesterol After Dropping Outliers')
plt.show()

<Axes: ylabel='totChol'>
Text(0.5, 1.0, 'Boxplot of Total Cholesterol After Dropping Outliers')

```



```
df_clean = df
```

```

scaler = MinMaxScaler(feature_range=(0,1))

#assign scaler to column:
df_scaled = pd.DataFrame(scaler.fit_transform(df_clean),
columns=df_clean.columns)

df_scaled.describe()
df.describe()

      sysBP      glucose       age     totChol   cigsPerDay
\count  3749.000000  3749.000000  3749.000000  3749.000000  3749.000000
mean    0.230991    0.118260    0.462432    0.352447    0.128728
std     0.104228    0.067429    0.225589    0.124179    0.170391
min     0.000000    0.000000    0.000000    0.000000    0.000000
25%     0.158392    0.087571    0.263158    0.264957    0.000000
50%     0.210402    0.107345    0.447368    0.344729    0.000000
75%     0.283688    0.132768    0.631579    0.430199    0.285714
max     1.000000    1.000000    1.000000    1.000000    1.000000

      diaBP  prevalentHyp      diabetes      BPMeds
male \
\count  3749.000000  3749.000000  3749.000000  3749.000000
3749.000000
mean    0.369677    0.311550    0.026941    0.030408
0.445185
std     0.126290    0.463189    0.161931    0.171730
0.497053
min     0.000000    0.000000    0.000000    0.000000
0.000000
25%     0.285714    0.000000    0.000000    0.000000
0.000000
50%     0.359788    0.000000    0.000000    0.000000
0.000000
75%     0.444444    1.000000    0.000000    0.000000
1.000000
max     1.000000    1.000000    1.000000    1.000000
1.000000

      TenYearCHD
count  3749.000000
mean    0.152307
std     0.359366

```

	sysBP	glucose	age	totChol	cigsPerDay
\count	3749.000000	3749.000000	3749.000000	3749.000000	3749.000000
mean	132.354628	81.863964	49.572419	236.708722	9.010936
std	22.044275	23.869703	8.572367	43.586786	11.927370
min	83.500000	40.000000	32.000000	113.000000	0.000000
25%	117.000000	71.000000	42.000000	206.000000	0.000000
50%	128.000000	78.000000	49.000000	234.000000	0.000000
75%	143.500000	87.000000	56.000000	264.000000	20.000000
max	295.000000	394.000000	70.000000	464.000000	70.000000
	diaBP	prevalentHyp	diabetes	BPMeds	
male \count	3749.000000	3749.000000	3749.000000	3749.000000	3749.000000
mean	82.934516	0.311550	0.026941	0.030408	0.445185
std	11.934410	0.463189	0.161931	0.171730	0.497053
min	48.000000	0.000000	0.000000	0.000000	0.000000
25%	75.000000	0.000000	0.000000	0.000000	0.000000
50%	82.000000	0.000000	0.000000	0.000000	0.000000
75%	90.000000	1.000000	0.000000	0.000000	1.000000
max	142.500000	1.000000	1.000000	1.000000	1.000000
	TenYearCHD				
count	3749.000000				
mean	0.152307				
std	0.359366				
min	0.000000				
25%	0.000000				

```
50%      0.000000
75%      0.000000
max      1.000000
```

#### Test - Train Split

```
# clarify what is y and what is x label
y = df_scaled['TenYearCHD']
X = df_scaled.drop(['TenYearCHD'], axis = 1)

# divide train test: 80 % - 20 %
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.2, random_state=29)

len(X_train)
len(X_test)

2999
750
```

#### Resampling imbalanced Dataset

```
import seaborn as sns
import matplotlib.pyplot as plt

# Checking balance of outcome variable
target_count = df_scaled.TenYearCHD.value_counts()
print('Class 0:', target_count[0])
print('Class 1:', target_count[1])
print('Proportion:', round(target_count[0] / target_count[1], 2), ': 1')

# Create the countplot with the x argument explicitly set
sns.countplot(x='TenYearCHD', data=df_scaled, palette="OrRd",
hue=None)


# Customize the plot
plt.box(False)
plt.xlabel('Heart Disease No/Yes', fontsize=11)
plt.ylabel('Patient Count', fontsize=11)
plt.title('Count Outcome Heart Disease\n')
plt.savefig('Balance_Heart_Disease.png')
plt.show()

Class 0: 3178
Class 1: 571
Proportion: 5.57 : 1
```

```
C:\Users\PC-6\AppData\Local\Temp\ipykernel_1176\3551037873.py:11:  
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be  
removed in v0.14.0. Assign the `x` variable to `hue` and set  
'legend=False' for the same effect.
```

```
sns.countplot(x='TenYearCHD', data=df_scaled, palette="OrRd",  
hue=None)  
  
<Axes: xlabel='TenYearCHD', ylabel='count'>  
Text(0.5, 0, 'Heart Disease No/Yes')  
Text(0, 0.5, 'Patient Count')  
Text(0.5, 1.0, 'Count Outcome Heart Disease\n')
```



```
import seaborn as sns  
import matplotlib.pyplot as plt
```

```

# Shuffle df
shuffled_df = df_scaled.sample(frac=1, random_state=4)

# Put all the heart disease cases in a separate dataset
CHD_df = shuffled_df.loc[shuffled_df['TenYearCHD'] == 1]

# Randomly select observations from the non-heart disease (majority class)
non_CHD_df = shuffled_df.loc[shuffled_df['TenYearCHD'] == 0].sample(n=611, random_state=42)

# Concatenate both dataframes
normalized_df = pd.concat([CHD_df, non_CHD_df])

# Check new class counts
print(normalized_df.TenYearCHD.value_counts())

# Plot new count
sns.countplot(x='TenYearCHD', data=normalized_df, palette="OrRd",
hue=None)
plt.box(False)
plt.xlabel('Heart Disease No/Yes', fontsize=11)
plt.ylabel('Patient Count', fontsize=11)
plt.title('Count Outcome Heart Disease after Resampling\n')
# plt.savefig('Balance Heart Disease.png')
plt.show()

TenYearCHD
0.0    611
1.0    571
Name: count, dtype: int64

C:\Users\PC-6\AppData\Local\Temp\ipykernel_1176\2733195083.py:20:
FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

    sns.countplot(x='TenYearCHD', data=normalized_df, palette="OrRd",
hue=None)

<Axes: xlabel='TenYearCHD', ylabel='count'>
Text(0.5, 0, 'Heart Disease No/Yes')
Text(0, 0.5, 'Patient Count')
Text(0.5, 1.0, 'Count Outcome Heart Disease after Resampling\n')

```

## Count Outcome Heart Disease after Resampling



```
y_train = normalized_df['TenYearCHD']
X_train = normalized_df.drop('TenYearCHD', axis=1)

from sklearn.pipeline import Pipeline

classifiers =
[LogisticRegression(), SVC(), DecisionTreeClassifier(), KNeighborsClassifier(2)]

for classifier in classifiers:
    pipe = Pipeline(steps=[('classifier', classifier)])
    pipe.fit(X_train, y_train)
    print("The accuracy score of {} is: {:.2f}%".format(classifier,
(pipe.score(X_test, y_test)*100)))

Pipeline(steps=[('classifier', LogisticRegression())])
The accuracy score of LogisticRegression() is: 65.73%
Pipeline(steps=[('classifier', SVC())])
```

```

The accuracy score of SVC() is: 64.80%
Pipeline(steps=[('classifier', DecisionTreeClassifier())])
The accuracy score of DecisionTreeClassifier() is: 70.13%
Pipeline(steps=[('classifier', KNeighborsClassifier(n_neighbors=2))])
The accuracy score of KNeighborsClassifier(n_neighbors=2) is: 80.00%

```

### Modelling & Evaluation (without Pipeline)

```

# logistic regression again with the balanced dataset

normalized_df_reg = LogisticRegression().fit(X_train, y_train)

normalized_df_reg_pred = normalized_df_reg.predict(X_test)

# check accuracy: Accuracy: Overall, how often is the classifier correct? Accuracy = (True Pos + True Negative)/total
acc = accuracy_score(y_test, normalized_df_reg_pred)
print(f"The accuracy score for LogReg is: {round(acc,3)*100}%")

# f1 score: The F1 score can be interpreted as a weighted average of the precision and recall, where an F1 score reaches its best value at 1 and worst score at 0.
f1 = f1_score(y_test, normalized_df_reg_pred)
print(f"The f1 score for LogReg is: {round(f1,3)*100}%")

# Precision score: When it predicts yes, how often is it correct?
Precision=True Positive/predicted yes
precision = precision_score(y_test, normalized_df_reg_pred)
print(f"The precision score for LogReg is: {round(precision,3)*100}%")

# recall score: True Positive Rate(Sensitivity or Recall): When it's actually yes, how often does it predict yes? True Positive Rate = True Positive/actual yes
recall = recall_score(y_test, normalized_df_reg_pred)
print(f"The recall score for LogReg is: {round(recall,3)*100}%")

The accuracy score for LogReg is: 65.7%
The f1 score for LogReg is: 37.2%
The precision score for LogReg is: 26.0%
The recall score for LogReg is: 65.0%


from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Ensure normalized_df_reg_pred is the prediction array

```

```
# cnf_matrix_log = confusion_matrix(y_test, normalized_df_reg_pred) #  
# Assuming it's already computed  
  
# Calculate confusion matrix  
cnf_matrix_log = confusion_matrix(y_test, normalized_df_reg_pred)  
  
# Plotting the heatmap  
sns.heatmap(pd.DataFrame(cnf_matrix_log), annot=True, cmap="Reds",  
fmt='g', cbar=False)  
  
# Customizing the axes and labels  
plt.title('Confusion Matrix Logistic Regression\n', y=1.1)  
plt.xlabel('Predicted Labels')  
plt.ylabel('True Labels')  
plt.tight_layout()  
  
# Show the plot  
plt.show()  
  
<Axes: title={'center': 'Confusion matrix Logistic Regression\n'}>  
Text(0.5, 1.1, 'Confusion Matrix Logistic Regression\n')  
Text(0.5, 9.444444444444459, 'Predicted Labels')  
Text(9.444444444444452, 0.5, 'True Labels')
```

Confusion Matrix Logistic Regression



```
# Support Vector Machine

#initialize model
svm = SVC()

#fit model
svm.fit(X_train, y_train)

normalized_df_svm_pred = svm.predict(X_test)

# check accuracy: Accuracy: Overall, how often is the classifier correct? Accuracy = (True Pos + True Negative)/total
acc = accuracy_score(y_test, normalized_df_svm_pred)
print(f"The accuracy score for SVM is: {round(acc,3)*100}%")

# f1 score: The F1 score can be interpreted as a weighted average of the precision and recall, where an F1 score reaches its best value at 1 and worst score at 0.
f1 = f1_score(y_test, normalized_df_svm_pred)
print(f"The f1 score for SVM is: {round(f1,3)*100}%")

# Precision score: When it predicts yes, how often is it correct?
```

```

Precision=True Positive/predicted yes
precision = precision_score(y_test, normalized_df_svm_pred)
print(f"The precision score for SVM is: {round(precision,3)*100}%")

# recall score: True Positive Rate(Sensitivity or Recall): When it's
# actually yes, how often does it predict yes? True Positive Rate = True
# Positive/actual yes
recall = recall_score(y_test, normalized_df_svm_pred)
print(f"The recall score for SVM is: {round(recall,3)*100}%")

SVC()

The accuracy score for SVM is: 64.8%
The f1 score for SVM is: 37.7%
The precision score for SVM is: 26.1%
The recall score for SVM is: 68.4%

import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from sklearn.metrics import confusion_matrix

# Assuming y_test and normalized_df_svm_pred are already defined

cnf_matrix_svm = confusion_matrix(y_test, normalized_df_svm_pred)

# Normalize confusion matrix if you want
# cnf_matrix_svm = cnf_matrix_svm.astype('float') /
# cnf_matrix_svm.sum(axis=1)[:, np.newaxis]

# Create DataFrame from confusion matrix
cnf_matrix_df = pd.DataFrame(cnf_matrix_svm,
                             index=np.unique(y_test),
                             columns=np.unique(y_test))

# Plot the heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(cnf_matrix_df, annot=True, cmap="Reds", fmt='g',
            cbar=True)

# Labels and title
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix - SVM', y=1.1)

# Adjust position of x-axis label
ax = plt.gca()
ax.xaxis.set_label_position("top")

# Tight layout to avoid clipping

```

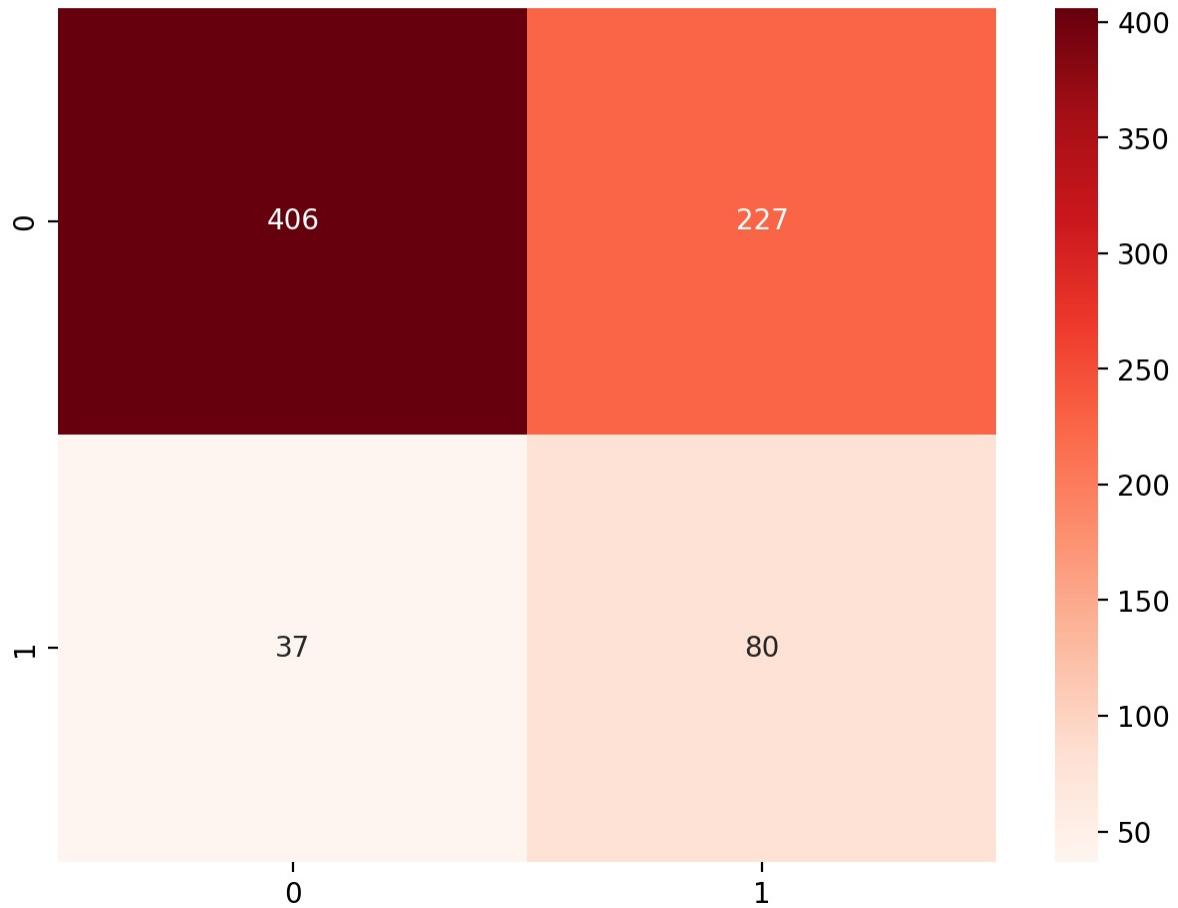
```
plt.tight_layout()
plt.show()

<Figure size 800x600 with 0 Axes>


<Axes: >

Text(0.5, 36.72222222222221, 'Predicted')
Text(70.72222222222221, 0.5, 'True')
Text(0.5, 1.1, 'Confusion Matrix - SVM')
```

Confusion matrix SVM



Confusion Matrix - SVM



```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score,
recall_score, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Initialize the Decision Tree Classifier model
dtc_up = DecisionTreeClassifier()

# Fit model
dtc_up.fit(X_train, y_train)

# Predict on test set
normalized_df_dtc_pred = dtc_up.predict(X_test)

# Accuracy: Overall, how often is the classifier correct?
acc = accuracy_score(y_test, normalized_df_dtc_pred)
print(f"The accuracy score for DTC is: {round(acc, 3) * 100}%")
```

```

# F1 score: The F1 score is the weighted average of precision and recall
f1 = f1_score(y_test, normalized_df_dtc_pred)
print(f"The F1 score for DTC is: {round(f1, 3) * 100}%")

# Precision score: When it predicts yes, how often is it correct?
precision = precision_score(y_test, normalized_df_dtc_pred)
print(f"The precision score for DTC is: {round(precision, 3) * 100}%")

# Recall score: When it's actually yes, how often does it predict yes?
recall = recall_score(y_test, normalized_df_dtc_pred)
print(f"The recall score for DTC is: {round(recall, 3) * 100}%")

# Confusion Matrix for more insights
cnf_matrix_dtc = confusion_matrix(y_test, normalized_df_dtc_pred)
cnf_matrix_df = pd.DataFrame(cnf_matrix_dtc,
                             index=np.unique(y_test),
                             columns=np.unique(y_test))

# Plotting the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cnf_matrix_df, annot=True, cmap="Blues", fmt='g',
cbar=True)

# Adding labels and title to the plot
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix - Decision Tree', y=1.1)

# Adjust position of x-axis label
ax = plt.gca()
ax.xaxis.set_label_position("top")

# Tight layout to avoid clipping
plt.tight_layout()
plt.show()

DecisionTreeClassifier()

The accuracy score for DTC is: 70.5%
The F1 score for DTC is: 51.4%
The precision score for DTC is: 34.59999999999994%
The recall score for DTC is: 100.0%

<Figure size 800x600 with 0 Axes>


<Axes: >

Text(0.5, 36.7222222222221, 'Predicted')

Text(70.7222222222221, 0.5, 'True')

```

```
Text(0.5, 1.1, 'Confusion Matrix - Decision Tree')
```



```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from sklearn.metrics import confusion_matrix

# Assuming y_test and normalized_df_dtc_pred are defined

# Compute confusion matrix
cnf_matrix_dtc = confusion_matrix(y_test, normalized_df_dtc_pred)

# Create DataFrame from confusion matrix for better clarity
cnf_matrix_df = pd.DataFrame(cnf_matrix_dtc,
                             index=np.unique(y_test),
                             columns=np.unique(y_test))

# Plot the confusion matrix
```

```
plt.figure(figsize=(8, 6))
sns.heatmap(cnf_matrix_df, annot=True, cmap="Reds", fmt='g',
cbar=True)

# Add axis labels
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix - Decision Tree', y=1.1)

# Adjust position of x-axis label
ax = plt.gca()
ax.xaxis.set_label_position("top")

# Tight layout to avoid clipping
plt.tight_layout()
plt.show()

<Figure size 800x600 with 0 Axes>


<Axes: >

Text(0.5, 36.7222222222221, 'Predicted')


Text(70.7222222222221, 0.5, 'True')

Text(0.5, 1.1, 'Confusion Matrix - Decision Tree')
```

Confusion matrix Decision Tree



Confusion Matrix - Decision Tree



```
from sklearn.model_selection import GridSearchCV

# Define the parameter grid for n_neighbors and metric
param_grid = {
    'n_neighbors': [1, 2, 3, 5, 7, 10],
    'metric': ['euclidean', 'manhattan', 'minkowski']
}

# Create a KNN classifier
knn = KNeighborsClassifier()

# Initialize GridSearchCV with 5-fold cross-validation
grid_search = GridSearchCV(knn, param_grid, cv=5, scoring='accuracy')

# Fit the model on training data
grid_search.fit(X_train, y_train)

# Get the best parameters from grid search
print("Best parameters from grid search:", grid_search.best_params_)
```

```

# Train the KNN model using the best found parameters
best_knn = grid_search.best_estimator_

# Prediction on test data
normalized_df_knn_pred = best_knn.predict(X_test)

# Evaluate the model

GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
             param_grid={'metric': ['euclidean', 'manhattan',
             'minkowski'],
                         'n_neighbors': [1, 2, 3, 5, 7, 10]},
             scoring='accuracy')

Best parameters from grid search: {'metric': 'euclidean',
'n_neighbors': 10}

# Assuming you have already defined best_knn after grid search
acc_test = best_knn.score(X_test, y_test)
print(f"The accuracy score of the test data is: {acc_test*100:.2f}%")

acc_train = best_knn.score(X_train, y_train)
print(f"The accuracy score of the training data is:
{acc_train*100:.2f}%")

The accuracy score of the test data is: 69.60%
The accuracy score of the training data is: 71.15%

import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.metrics import confusion_matrix

# Assuming y_test and normalized_df_knn_pred are defined
cnf_matrix_knn = confusion_matrix(y_test, normalized_df_knn_pred)

# Create the heatmap
plt.figure(figsize=(8, 6)) # Optional: adjust the figure size for clarity
ax = plt.subplot()
sns.heatmap(pd.DataFrame(cnf_matrix_knn), annot=True, cmap="Reds",
fmt='g', cbar=False,
           xticklabels=['Class 0', 'Class 1'], yticklabels=['Class 0', 'Class 1'])

# Set labels and title
ax.set_xlabel('Predicted')
ax.set_ylabel('True')
ax.set_title('Confusion Matrix for KNN')

```

```

# Show the plot
plt.show()

<Figure size 800x600 with 0 Axes>


<Axes: >

Text(0.5, 36.72222222222221, 'Predicted')

Text(70.72222222222221, 0.5, 'True')

Text(0.5, 1.0, 'Confusion Matrix for KNN')

```



```

# AU ROC CURVE KNN
'''the AUC ROC Curve is a measure of performance based on plotting the
true positive and false positive rate
and calculating the area under that curve.The closer the score to 1
the better the algorithm's ability to
distinguish between the two outcome classes.''''

```

```

fpr, tpr, _ = roc_curve(y_test, normalized_df_knn_pred)
auc = roc_auc_score(y_test, normalized_df_knn_pred)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.box(False)
plt.title ('ROC CURVE KNN')
plt.show()

print(f"The score for the AUC ROC Curve is: {round(auc,3)*100}%")


"the AUC ROC Curve is a measure of performance based on plotting the
true positive and false positive rate \nand calculating the area under
that curve.The closer the score to 1 the better the algorithm's
ability to \ndistinguish between the two outcome classes."

```

[<matplotlib.lines.Line2D at 0x2394db81eb0>]

<matplotlib.legend.Legend at 0x2394b1b9100>

Text(0.5, 1.0, 'ROC CURVE KNN')



The score for the AUC ROC Curve is: 63.5%

```

def start_questionnaire():
    my_predictors = []
    parameters=['sysBP',
    'glucose','age','totChol','cigsPerDay','diaBP','prevalentHyp','diabetes','BPMeds','male']

    print('Input Patient Information:')

    age = input("Patient's age: >>> ")
    my_predictors.append(age)
    male = input("Patient's gender. male=1, female=0: >>> ")
    my_predictors.append(male)
    cigsPerDay = input("Patient's smoked cigarettes per day: >>> ")
    my_predictors.append(cigsPerDay)
    sysBP = input("Patient's systolic blood pressure: >>> ")
    my_predictors.append(sysBP)
    diaBP = input("Patient's diastolic blood pressure: >>> ")
    my_predictors.append(diaBP)
    totChol = input("Patient's cholesterol level: >>> ")
    my_predictors.append(totChol)
    prevalentHyp = input("Was Patient hypertensive? Yes=1, No=0 >>> ")

    my_predictors.append(prevalentHyp)
    diabetes = input("Did Patient have diabetes? Yes=1, No=0 >>> ")
    my_predictors.append(diabetes)
    glucose = input("What is the Patient's glucose level? >>> ")
    my_predictors.append(glucose)
    BPMeds = input("Has Patient been on Blood Pressure Medication? Yes=1, No=0 >>> ")
    my_predictors.append(BPMeds)

    my_data = dict(zip(parameters, my_predictors))
    my_df = pd.DataFrame(my_data, index=[0])
    scaler = MinMaxScaler(feature_range=(0,1))

    # assign scaler to column:
    my_df_scaled = pd.DataFrame(scaler.fit_transform(my_df),
columns=my_df.columns)
    my_y_pred = knn.predict(my_df)
    print('\n')
    print('Result:')
    if my_y_pred == 1:
        print("The patient will develop a Heart Disease.")
    if my_y_pred == 0:
        print("The patient will not develop a Heart Disease.")

start_questionnaire()
Input Patient Information:

```

